

ThinLisp Manual

Version 0.5
10 October 1999

The ThinLisp Group:
Jim Allard
Ben Hyde

Copyright © 1999 The ThinLisp Group
Copyright © 1995 Gensym Corporation.
All rights reserved.

This file is part of ThinLisp.

ThinLisp is open source; you can redistribute it and/or modify it under the terms of the ThinLisp License as published by the ThinLisp Group; either version 1 or (at your option) any later version.

ThinLisp is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

For additional information see <http://www.thinlisp.org/> .

Preface

ThinLisp is an open source Lisp to C translator for delivering commercial quality, Lisp-based application. It implements a subset of Common Lisp with extensions. ThinLisp itself is written in Common Lisp, and so must run on top of an underlying Common Lisp implementation such as Allegro, MCL, or CMU Lisp. The C code resulting from a translation can then be independently compiled to produce a small, efficient executable image.

Originally designed for real-time control applications, ThinLisp is not a traditional Lisp environment. It purposefully does not contain a garbage collector. ThinLisp produces compile time warnings for uses of inherently slow Lisp operations, for consing operations, and for code that is unoptimizable due to a lack of sufficient type declarations. These warnings can be suppressed by improving the code, or through use of lexical declarations acknowledging that the code is only of prototype quality. Code meeting the stringent requirements imposed by ThinLisp cannot be sped up by rewriting it in C.

The ThinLisp home is <http://www.thinlisp.org/>. The newest source distributions of TL and this manual can be found here. Bugs should be reported to bugs@thinlisp.org.

Acknowledgements

ThinLisp was begun in 1995 as a late night project born of the frustration of attempting to ship commercial quality software products written in Lisp. Four years later I'm still obsessed by the issues that drove me to it in the first uplace — practical use of high level languages to implement complex applications while achieving the performance that seems easy when using languages that are semantically "close to the silicon". Hopefully this implementation will lead to some useful insights towards that goal.

Thanks to Ben Hyde for convincing Gensym and myself that this system should be taken off the shelf, polished up, and released as open source software. Mike Colena and Glen Iba wrote significant parts of the implementation. Nick Caruso, Joe Devlin, Rick Harris, and John Hodgkinson and also contributed their time, sage opinions, and code. Thanks to Lowell Hawkinson, Jim Pepe, and Dave Riddell for their support of this work at Gensym Corporation. Thanks to Kim Barrett and David Sotkowitz of IS Robotics, Inc. for contributions of multiple-inheritance code and opinions, and to Rod Brooks and Dave Barrett for their ideas. Finally, thanks to my wife, Gerry Zipser, for putting up with the bleary-eyed aftermath of coding sessions, and for helping to focus efforts during this past summer's release push.

Jim Allard
Newton, Massachusetts
October 6, 1999

1 Rant

1.1 My Excuse for Ranting

I've imagined the reactions people will have when they see that there is yet another Lisp to C translator implementation being released. Mostly I've expected that people would think I'm stark raving mad to engage in a large work, to serve an almost non-existent user base of programmers in a language that is largely discredited while being actively antagonistic towards many of the design principles this language has employed for decades. Since I'm expecting that people will think I'm mad (and half believing it myself), I then feel entitled (even obligated) to engage in a rant to explain the motivations of such a curious endeavor. Please excuse this indulgence and look charitably upon it as a form of self-therapy.

1.2 Lisp Background

More than fifteen years ago I was in a recitation session of an Introduction to Artificial Intelligence course. The lecturer, Patrick Winston, had come in for a session on automated planning. At one point I asked a question about how a certain situation might be handled in the system we were studying, and he answered that it couldn't be handled in that system and that solving that problem was a research topic. The fact that problems as simple (so I thought at the time) as what we were describing were research, and that as a sophomore in college I could grasp them seemed amazing to me then, as it still does now. On the spot I could imagine six different ways to attack that problem, and none of them felt any less plausible than some of the established techniques that we were studying. I was hooked.

At the time, this sort of optimism and arrogance about the immediacy and inevitability of technical success in AI was rampant. The convergence of breakthroughs in techniques, tools, and venture capital money that flowed like wine was intoxicating for everyone.

No one thought this work would be trivial, so the best software development tools available would be required. And of course the best tools from the perspective of the AI community included the language that this community had been developing since the late 1950's, Lisp.

Originally, Lisp was developed as an experiment in providing a direct means of implementing lambda calculus, where operations were first class objects that could be manipulated in the same manner as data. The presumption behind many of the Lisp design choices was that the most important goal was to maximize a programmer's productivity in developing new functionality by enabling the purest form possible of coded solutions, and so to hide as many of the platform specific details as possible. In the ongoing evolution of Lisp, considerations of the efficiency of processing speed and memory use came in a distant second to programmer convenience and to the attempt to protect the programmer from bugs caused by mathematically abhorrent behaviors such as integer wraparound.

Perhaps because of this devotion to developer productivity; perhaps because of the synergy between language syntax, parsing, and data; perhaps because of the invention of the Lisp macro facility; or perhaps because it was invented down the hall, Lisp became the

language of choice for research and experimentation in new programming constructs (at least at MIT). Many of the different possible approaches to object oriented programming were first tried out in Lisp. Examples are multiple inheritance, multi-methods, prototype-based object systems, and automatic method combination via whoppers, wrappers, before, and after methods. Other innovations experimented with in Lisp are error handling systems, multi-platform file naming approaches, and guaranteed execution of clean-up code (i.e. `unwind-protect`).

In 1984 a group formed to attempt to unify as many of the best features of different Lisp dialects as possible into one central standard. From this effort came the excellent books *Common Lisp, The Language* by Guy Steele and its second edition, which incorporated many of the changes that would eventually be codified into ISO Common Lisp. This is the dialect that dominates Lisp work today, at least in the U.S.

1.3 Mathematician's Languages

From the beginning Lisp was a mathematician's language. To understand what I mean by Lisp this, I must explain my view of the differences between the goals of scientists, mathematicians, and engineers.¹

The goal of scientists is to learn truths about the world around them, using observation, analysis, hypothesis, and experimentation to advance and buttress their claims of discovery. Pure science aspires to being unconcerned with the utility or applicability of the information they seek. The ultimate goal is to discover what *is*.

Mathematicians aspire to an even higher goal — understanding what truths there are or that might be, regardless of or even entirely divorced from the limits of the physical world. The studies of mathematicians are the stuff of pure thought and imagination, yet bound by the strictest of rules of proof. To build robust castles in the air based on proofs combining into a single monolithic whole requires that all pieces must act perfectly in concert. Therefore the mathematician must constantly be vigilant against contradictions between the different portions of his whole proof, where one technique carries with it constraints which must be respected within all other parts of the proof. With this rigorous technique in hand, the ultimate goal is to discover what is *true*.

On the other hand the goal of the engineer, in the abstract, is to bend the behavior of an uncooperative world to his will. Whether it is the stereotypical train engineer, civil engineer, acoustic, electrical, mechanical, or software engineer, their work is to use the fruits of science and mathematics, exploiting practicality and expediency, to operate or design a physical entity so that it performs a specific task, usually with the motivation of the almighty dollar. Within the constraints of the physical world, the ultimate goal is to accomplish what is *practical*.

Most computer languages in use today are exemplars of the engineering style. Building primarily from the operations that could be implemented on the computing devices that could be made at the time, and perhaps secondarily on the needs of a user base at hand, languages were designed to allow people to command computers to do things. Typically the

¹ All apologies in advance to my father-in-law (a scientist) and my brother-in-law (a mathematician).

operations in the language were semantically thin (or shallow, if you prefer), to the extent that a programmer of medium experience could almost predict the number of machine instructions required to carry out any given command in the language. Fortran, Pascal, Basic, and C are all examples of languages in this category, with C (and its derivative, C++) now being the default choice for people building big systems.

Lisp and other Languages (such as Smalltalk, Eiffel, and recently Java) were developed along a different tack. The operations in these languages were defined with deep semantics in mind. Here the focus was not on what an operation *did* but on what it *meant*. Of course there had to be implementations of these operations, but the actual machine instructions carried out were allowed (even expected) to vary widely depending on the manner and situation in which it was used.

An example of how this difference in design perspective plays out can be found in the behavior of the `goto` statement in C and the `go` special operator in Common Lisp. Consider the following two code fragments.

In C:

```
int find_big_int(int *set, int start, int end) {
    Iterate through the set {
        if test
            goto bad;
        ...
    }
    bad:
    ...
}
```

In Lisp:

```
(defun find-big-int (set start end)
  (tagbody
    (Iterate through the set
      (if test
        (go bad))
      ...))
    bad
    ...))
```

Describe why a lexical closure surrounding the `(go bad)` could make this very expensive.²

1.4 Lisp in Action

In any case, the bulk of the research labs and companies that attempted to commercialize AI technologies used Common Lisp to implement their products. The most broadly

² This needs further work. Describe Godelizing, the tendency of "mathematical" programmers to choose such correct, but overweight implementation approaches, and the underestimation by the abstract minded of the difficulty of abstracts to the teams of people needed to finish a complex software system.

deployed AI-based technique was expert systems, also called rule-based systems.³ Early results were encouraging. In the lab, sometimes on specialized hardware, the software engineers were able to demonstrate the beginnings of the functionality that they were promising. However, when it came to deploying these systems, things didn't turn out so smoothly. Most software products of that era were coded in languages such as Fortran, Cobol, Pascal, C, and an assembly language. It turned out that Software developers who could breathe fire in Lisp were clumsy, inefficient, or flatly unwilling to port their systems to these traditional languages.

which were used at the time to deliver the majority of commercial products

Customers were loathe to buy expensive, special purpose hardware for these applications while cheap, ever more powerful PC computers were rolling out. Stories of this sort are rife throughout the industry, but I'll speak to just the cases I personally was involved in.

³ Of course this is ignoring other, more broadly applicable developments, such as multi-tasking operating systems.

2 Original Introduction

This introduction was written at the beginning of the effort to make a new Lisp to C translator. At the time I was working for Gensym Corporation, and named the project after the company. -jallard 10/9/99

The Gensym Language Translator (GLT) is being made to enable Gensym to control the means of distribution of its products. Since 1986 till now (January, 1995), we have used 6 different brands of Lisp compiler, abandoning each one in turn for platform availability, portability, and performance reasons. Most recently we have been using the Chestnut Lisp to C translator to deliver our Lisp-based products for releases 3 and 4. Between these two releases of our products, we took a new release from Chestnut, and it took a full man-year of my time to resolve the problems that had been introduced. I've made the judgment that for less labor than is devoted to maintaining our use of Chestnut's product, we can implement and maintain our own translator. It's my intent to implement a translator that meets or beats Chestnut's characteristics for the portions of Lisp that we use, and that does not waste time or attention on the portions that we do not use. Time can be taken later to extend it into areas we would like to use, but currently cannot because we can't depend on non-consing implementations.

The base of GLT consists of a small hand-written C include file, a small hand-written memory allocation C file, and a translator written in Common Lisp that can translate a minimal set of primitives from Lisp to C. The primitives in the translator implement data structure allocation and manipulation, the defining of constants, parameters, variables, macros, and functions. The system will not contain many of the features normally present in a full Common Lisp implementation, most notably no garbage collector, lexical closures, rest arguments, condition system, bignums, complex numbers, or CLOS.¹ The built-in types will be limited to packages, symbols, conses, fixnums, doubles, simple-vectors, simple-arrays of unsigned-byte 1, 8, and 16 bits, simple-arrays of doubles, and fill-pointered non-adjustable strings. Notable complex features that will be included are special variables, catch and throw, unwind-protect, multiple-values, optional arguments, defstruct, macrolet, labels, flet, and deftype.

Particular attention will be paid to using C integer and double types for variables, function arguments, and function values when type declarations allow, largely because it bugs me when these things aren't handled right and I believe it is a significant performance loss to always be boxing and unboxing at function boundaries. We will also make special efforts to optimize calls to functions that return a single value.

¹ Rest arguments have since been added. -jallard 10/9/99.

3 Memory Architecture

Use pointers to a one word (4 byte) header structure as the basic mechanism. The type should be an 8-bit unsigned byte value. The upper 3 bytes of the header word should be used for type specific data. For vectors this should be the length, giving us a simple vector with one word of overhead (all previous Lisp implementations we've used, except the LispM, had two or more words of header). All Lisp objects should be aligned on 4 byte addresses. There is an argument that they should be aligned on 8 byte addresses. There are two advantages to 8 byte alignment. One is that it would enable us to reliably store immediate double floats and pointers into the same array. The second is that it would can give us an immediate type tag for conses that is itself the offset to the cdr of a cons. The argument for 4 byte alignment is that all current C compilers we use align structures on 4 byte addresses, but some won't align them on 8 byte addresses, even if they contain doubles. The other argument for 4 byte alignment is that we would not need to occasionally skip forward 4 bytes when allocating from heap in order to find the next 8 byte aligned address. For now, the 4 byte alignment wins.

The exceptions to the pointer to a type tag rule are the following types, which have immediate type tags. These types are fixnums, conses, and managed-floats. If all pointers are 4 byte aligned, then the lower 2 bits of all pointers are always zeros. This gives us 3 non-zero immediate type tags. The following are the immediate type tag assignments:

- 0: pointer to header of Lisp object
- 1: immediate fixnum
- 2: pointer to cons (mask with -4 for car*, follow for cdr*)
- 3: immediate character

A determination of type can be made from a hex value of the pointer.

- hex 0, 4, 8, C are pointers to Lisp objects
- hex 1, 5, 9, and D are fixnums,
- hex 2, 6, A, and E are pointers to conses
- hex 3, 7, B, and F are immediate characters.

The arguments for a 30 bit fixnum instead of a 31 is that it makes for fast fixnum additions and subtractions without risking overflow. 31 bit fixnums require a sequence point between operations to avoid the potential for overflow.

Characters could be either immediate or remote, though there is some performance benefit to characters being immediate. We could have a preallocated array of remote character objects, and allocate a Lisp character is arefing into this array. For now we are going with immediate characters, especially since we are thinking hard about UNICODE characters, which require 16 bits and so make a pre-allocated array become somewhat too large.

The type tags in headers should not conflict with any of the immediate type tags, so that type-case can turn into a fixnum-case of the type values.

Given these set ups, type tests against straight types are at worst a null test, a mask and an integer equality test, a character byte fetch, and another integer equality test.

In C, the type for object should be an unsigned integer type 32 bits long. There are several advantages to this. All architectures (including the Alpha OSF when using a linker

option) can have pointer values be represented in 32 bits, but for some platforms (the Alpha OSF) pointer types can consume something other than 4 bytes, 8 for all pointers on the Alpha OSF. By forcing all objects to consume 4 bytes, we can get interesting packing in structures and vectors on all platforms. [The following comment applies only to 8 byte alignment: For example, we could put immediate floats into odd-indexed simple vector locations, and those floats would consume 2 elements. For structures and frames, this could provide significant savings. -jra 8/30/95]

The type for managed float should have an immediate type tag so that we may have the smallest possible representation for floats, since so many are used. For each type that involves a heap allocated block of memory (i.e. all but fixnum and characters), there should be a corresponding type for a reclaimed instance of that type. This gives us a fast means of testing if a data structure is currently reclaimed or not, and would cause type tests in a safe translation. Even in production systems, we could check for double reclamation of data structures. Note that this is not possible for heap allocated data structures that use an immediate type tag.

The type tags for heap allocated data structures have no special issues, except that these values should not collide with any immediate type tags (i.e. be greater than 3), and be able to quickly determine if a data structure is reclaimed (i.e. use bit 7 as a flag).

So the type tag table is as follows:

Tag	Value	Reclaimed	Tag	
Dec	Hex	Lisp	Type C	Type & Conv Dec Hex

0 00		immed pointer	(Header *)Obj	
1 01		immed fixnum	((sint32)Obj)>>2	
2 02		cons	(Obj *) (Obj-2)	
3 03		immed character	(unsigned char)(Obj>>2)	
4 04		managed-float	(Mdouble *)Obj	132 84
5 05		double-float	(Ldouble *)Obj	133 85
6 06		simple-vector	(Sv *)Obj	134 86
7 07		string (w/fill ptr)	(Str *)Obj	135 87
8 08		(simple-array ubyte 8)	(Sa_uint8 *)Obj	136 88
9 09		(simple-array ubyte 16)	(Sa_uint16 *)Obj	137 89
10 0A		(simple-array double)	(Sa_double *)Obj	138 8A
11 0B		symbol	(Sym *)Obj	139 8B
12 0C		compiled-function	(Func *)Obj	140 8C
13 0D		package	(Package *)Obj	141 8D

Each of the data structures are described below.

3.1 Simple-vectors

Simple vectors are the most often used data structure, so we'd like to keep it as small and fast as possible. It has a one word (i.e. 4 byte) header. The only components of a simple vector are the type tag, the length, and the body of the array. If the type tag is 8 bits wide (unsigned) then the length can be 24 bits of unsigned integer, giving a maximum length of 16 Meg.

```
typedef struct {
    unsigned int:8:type;
    unsigned int:24:length;
    Obj[1]:body;
} Sv
```

In C references to array are guaranteed to not be bounds checked. This means that this one type can be used to all elements of arbitrarily sized simple vectors. Constant vectors can be made by having a type of simple-vector local to the C file containing the constant so that we can use an initialized structure. For example, a constant simple vector containing 5 fixnums could be emitted as C code as follows.

```
typedef struct {
    unsigned int:8:type;
    unsigned int:24:length;
    Obj[5]:body;
} Sv_5

static Sa_5 const1 = {8, 5, {fix(23), fix(2), fix(3),
    fix(9), fix(-12)}};
```

An example of the translation of a fetch of a simple vector element follows.

```
Lisp: (setq x (svref y 5))
C:    x = (Sv *)y->body[5];
```

Though the detail of casting and fetching the body component could be hidden in a C macro, at first we will leave all the details exploded out. One criticism of Chestnut's translations is that no one can figure out what the implementation is actually doing. Exploding out the details will help train development in the details.

3.2 Strings

In order to give fast performance for all strings, we will have fill pointers for all strings in G2. The fill pointer and length will both consume 3 bytes, plus one byte for the tag. The body of the string will be packed directly against the fill pointer, meaning that the most efficient packing of strings into our 8 byte aligned space will be for strings with lengths that have $\text{mod}(\text{length}, 8) = 1$. Since strings in C must be null terminated, that adds 1 extra byte to the length as compared to the loop length of the string. Having 1 extra byte in the header word will allow us to efficiently pack strings into words when they have lengths that are multiples of 4.

```
typedef struct {
    unsigned int: 8: type;
    unsigned int:24: length;
    unsigned int:24:fill_length;
    char[9]: body
} Str
```

3.3 Immediate integer and double arrays

All of these types are implemented in the same way, with a 4 byte-aligned one byte type tag, a 3 byte length, and then a body. These types are `Sa_uint8` and `Sa_uint16`. `Sa_double` is similar except that it is 8 byte aligned.. The type for bit-vectors is built on top of `Sa_uint8`, by fetching bytes and bit-shifting to fetch values and modify values.

3.4 Symbol

There are several different optimizations that one can imagine for the slots of symbols. In the first pass of this implementation, no such memory squeezing will be attempted, and all five typical slots of symbols will be directly provided in the symbol structure. Also, since virtually all symbols are included in packages in G2 and TW, the slots needed for holding the symbol in package balanced binary trees will be included in the symbol itself.

```
typedef struct {
    unsigned int:8:type;
    unsigned int:1:local_value;
    unsigned int:1:external;
    signed int:4:balance;
    unsigned int:1:imported;
    unsigned int:16:name_hash;
    Obj:symbol_name;
    Obj:symbol_value;
    Obj:symbol_plist;
    Obj:symbol_package;
    Obj:symbol_function;
    Obj:package_left_branch;
    Obj:package_right_branch;
} Symbol
```

The symbol-value slot is typically a pointer to the location containing the true symbol value, which will typically be a C global variable. When a runtime generated symbol has a value, the local-value bit will be 1, and the value is directly within the symbol-value slot. Note that slows down explicit calls to Lisp symbol-value since it has to check the local-value bit, but these explicit calls are rare and so I don't think that's so bad.

If packages are implemented as balanced binary trees, using the hash value of the name as an index, then there can be collisions between these hash values. In this case, the symbol names will be alphabetically ordered to determine left, right or match. These choices give us a constant size for symbols, given the size of the symbol-name.

3.5 Compiled Functions

Compiled functions will contain a pointer to the C function for this compiled function, the number of arguments for the function; the number of those arguments that are optional in the Lisp function, and a list of the constants that are default values for the optional arguments. Within the C runtime system, all functions will receive all arguments. If a funcall of a compiled function occurs where some optional arguments are not provided, the

default values are extracted from the list of default values. (In circumstances where the default value semantics require more than a constant, the no-arg argument value will be passed, and further computation will happen within the function to implement the default value selection).

```
typedef struct {
    unsigned int:8:type;
    unsigned int:8:arg_count;
    unsigned int:8:optional_arguments;
    Obj:default_arguments;
    void *f(): c_function;
} Func
```

The type spec for compiled functions will typically have arguments of type Obj: and a return type Obj: In special cases we may provide unsafe optimizations for functions that do not have object type arg and return values, but these should never be given to the default funcall operations.

3.6 Characters

Since all immediate type tags are in use, we will use heap allocated character objects. To avoid the need to reclaim them, we will preallocate all 256 character objects (at 8 bytes per, that's 2K). The value is explicitly unsigned to avoid confusion in parts that have a default signed char type.

```
typedef struct {
    double [0]: align;
    unsigned int:8:type;
    unsigned char:value
} Lchar
```

3.7 Package

Packages are implemented as balanced binary trees. They contain a list of used packages, and a symbol (or null pointer) that is the root node of the balanced binary tree. The balanced binary tree of symbol structures will hold only symbols that are local to this package and that have been imported to this package. External symbols of a package are represented by mask bits on the symbol itself. Imported symbols are represented by a new symbol structure in the binary tree, and its symbol value points to the imported symbol. The intern function will be implemented by searching through the binary trees of this package then through the binary trees of used packages and their used packages, but only accepting external symbols.

```
typedef struct {
    unsigned int:8:type;
    Obj:root_symbol;
    Obj:used_package_list;
} Pkg
```

3.8 Obsolete 8 Byte Alignments

[The following section describes tagging for 8 byte alignments. Keeping for historical perspective. -jra 8/30/95]

The exceptions to the pointer to a type tag rule are the following types, which have immediate type tags. These types are fixnums, characters, managed-floats, and conses. If all pointers are 8 byte aligned, then the lower 3 bits of all pointers are always zeros. This gives us 7 non-zero immediate type tags. The following are the immediate type tag assignments:

```
0: pointer to header of Lisp object
1: immediate even fixnum
2: immediate character
3: pointer to managed-float (mask with -8 for double*)
4: pointer to cons (mask with -8 for car*, follow for cdr*)
5: immediate odd fixnum
6: pointer to double-float (mask with -8 for double*)
7: unused
```

With this implementation, some determination of type can be made from the printed hex value of the pointer:

```
hex 0 and 8 are pointers to Lisp objects,
hex 1, 5, 9, and D are fixnums,
hex 2 and A are characters,
hex 3 and B are managed-floats,
hex 4 and C are conses,
hex 6 and E are double-floats,
hex 7 and F are unused, so must be corrupted objects
```

By having a cons tag of 4 and a headerless cons, the access to one of the elements of the cons will be a straight pointer dereference. The other requires subtracting 4 from the pointer first. (On the Alpha OSF or other 8 byte pointer, machines, some address computation is necessary for both the car and cdr, unless we make object a 4 byte wide data type, which we could do).

The type tags for heap allocated data structures have no special issues, except that these values should not collide with any immediate type tags (i.e. be greater than 7), and be able to quickly determine if a data structure is reclaimed (i.e. use bit 7 as a flag).

So the type tag table is as follows:

Tag	(Hex)	Lisp Type	C Type	Reclaimed Tag
0	00	immed	pointer (Header *)Obj	
1	01	immed	even fixnum (sint32)Obj>>2	
2	02	immed	character (char)(Obj>>2)	
3	03	managed-float	(double *)Obj	
4	04	cons	(Obj *)Obj-1	
5	05	immed	odd fixnum (sint32)Obj>>2	
6	06	double-float	(double *)Obj	
7	07	unused		
8	08	simple-vector	(Sv *)Obj	136 88
9	09	string (w/kill ptr)	(Str *)Obj	137 89
10	0A	(simple-array ubyte 8)	(Sa_uint8 *)Obj	138 8A
11	0B	(simple-array ubyte 16)	(Sa_uint16 *)Obj	139 8B
12	0C	(simple-array double)	(Sa_double *)Obj	140 8C
13	0D	symbol (Sym *)Obj		141 8D
14	0E	compiled-function (Func *)Obj		142 8E
15	0F	package (Package *)Obj		143 8F

Note the use of a zero length double array at the beginning of the struct. This ensures double word alignment, which gives us two things. The first is that pointers to constant structures of this type will be on 8 byte boundaries, required for our use of 3 bits of immediate type tag. The second benefit is that with a known alignment we can reliably mix objects and immediate doubles in simple-vectors. This is needed for structure and frame extensions I would like to make. The Harbison and Steele manual suggests that this zero length double array technique will ensure alignment, but this has not yet been tested.

Index

R

rant 3

Table of Contents

Preface	1
Acknowledgements	2
1 Rant	3
1.1 My Excuse for Ranting	3
1.2 Lisp Background	3
1.3 Mathematician's Languages	4
1.4 Lisp in Action	5
2 Original Introduction	7
3 Memory Architecture	8
3.1 Simple-vectors	9
3.2 Strings	10
3.3 Immediate integer and double arrays	11
3.4 Symbol	11
3.5 Compiled Functions	11
3.6 Characters	12
3.7 Package	12
3.8 Obsolete 8 Byte Alignments	13
Index	15